

Research Question 3

Observing Data

We perform a simple look at the means of the test score of each combination of TrDs. We can from this look at which combination of training sets achieves the maximum test score.


```
#library(emmeans)
#setwd("~/SeminarDataScience/semimar-data-sci/Coursework-B/RQ-3")
d = read.csv('../data/data.csv')
means = aggregate(score~TrD1+TrD2+TrD3+TrD4+TrD5+TrD6+TrD7+TrD8,d,mean)
head(means[order(-means$score),], n=25)
```

##	TrD1	TrD2	TrD3	TrD4	TrD5	TrD6	TrD7	TrD8	score
## 140	0	1	1	0	1	0	1	1	0.5660503
## 119	0	1	1	0	1	1	0	1	0.5624479
## 101	0	1	1	0	1	0	0	1	0.5618264
## 148	0	1	1	1	1	0	1	1	0.5605164
## 128	1	1	1	0	0	0	1	1	0.5561881
## 131	0	1	0	1	0	0	1	1	0.5524732
## 122	0	1	0	1	1	1	0	1	0.5498406
## 107	1	1	1	1	1	0	0	1	0.5495461
## 72	0	1	0	1	0	1	1	0	0.5490720
## 163	0	0	1	0	1	1	1	1	0.5489510
## 65	1	1	1	1	1	0	1	0	0.5485297
## 56	0	1	0	0	1	0	1	0	0.5464158
## 76	0	1	0	0	1	1	1	0	0.5462193
## 153	1	1	1	0	0	1	1	1	0.5454987
## 116	0	1	0	0	1	1	0	1	0.5453343
## 151	0	1	0	0	0	1	1	1	0.5450413
## 141	1	1	1	0	1	0	1	1	0.5435354
## 99	0	1	0	0	1	0	0	1	0.5427647
## 106	0	1	1	1	1	0	0	1	0.5415554
## 71	1	1	1	0	0	1	1	0	0.5411909
## 61	0	1	0	1	1	0	1	0	0.5406027
## 142	0	0	0	1	1	0	1	1	0.5401571
## 136	0	1	0	0	1	0	1	1	0.5397787
## 172	0	1	1	1	1	1	1	1	0.5397223
## 127	0	1	0	0	0	0	1	1	0.5394829

From this we can see that most of the combinations with the highest performance scores contain the Training Sets TrD 2 and 8. Further we can see that the top 3 combinations also contain TrD 3 and 5 potentially making these training sets also relevant. However this can also be a result of simply the use of more training sets improving the score.

Thus we want to analyse the effect of the amount of training sets used on the performance.

```
d$TrainCount <- apply(d[c(2:9)], 1, function(x) length(which(x==1)))
boxplot(score~TrainCount, d,xlab = "Amount of Training Sets", ylab = "Scores")
```



```
modelCount <- lm(score~TrainCount,d)
summary(modelCount)
```

```
##
## Call:
## lm(formula = score ~ TrainCount, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.50363 -0.18329  0.06978  0.16911  0.43422
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.442061   0.009987  44.264 < 2e-16 ***
## TrainCount   0.013036   0.002175   5.995 2.28e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2473 on 2994 degrees of freedom
## Multiple R-squared:  0.01186,    Adjusted R-squared:  0.01153
## F-statistic: 35.93 on 1 and 2994 DF,  p-value: 2.285e-09
```

```
confint(modelCount)
```

```
##              2.5 %      97.5 %
```

```
## (Intercept) 0.422478910 0.46164286
## TrainCount 0.008771912 0.01729971
```


Looking at the boxplot we can see that the amount of training sets used in general indeed has an effect on the score, although this effect seems to be minimal after 3 training sets used. Additionally also the linear model shows a sufficiently low p value, thus we can reject the null-hypothesis, meaning that there is a relevant relation between the score and the amount of training sets used. In this case the relation is one where the score increases with an increase in training sets.

We can also observe the single Training data set models and see how they perform

```
t1 = d[d$TrD1 == 1 & d$model=="S", ]
t2 = d[d$TrD2 == 1 & d$model=="S", ]
t3 = d[d$TrD3 == 1 & d$model=="S", ]
t4 = d[d$TrD4 == 1 & d$model=="S", ]
t5 = d[d$TrD5 == 1 & d$model=="S", ]
t6 = d[d$TrD6 == 1 & d$model=="S", ]
t7 = d[d$TrD7 == 1 & d$model=="S", ]
t8 = d[d$TrD8 == 1 & d$model=="S", ]

boxplot(t1$score,t2$score,t3$score,t4$score,t5$score,t6$score,t7$score,t8$score,names = c("1","2","3","4","5","6","7","8"))
```

Overall Performance of single training set models



From this analysis we can see that for the simple models, those trained on either TrD 5,7,8 have the best average general performance for all testdatasets. Important to note is that the scores when using TrD 2 show a observably larger variance when compared to the other sets.

Linear models

Now we make a linear model with the score and all the training data sets.

```
m <- lm(score~TrD1+TrD2+TrD3+TrD4+TrD5+TrD6+TrD7+TrD8, d)
summary(m)

##
## Call:
## lm(formula = score ~ TrD1 + TrD2 + TrD3 + TrD4 + TrD5 + TrD6 +
##     TrD7 + TrD8, data = d)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.51582 -0.18532  0.06918  0.17583  0.43854
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.439589   0.010191  43.133 < 2e-16 ***
## TrD1        -0.020523   0.009717  -2.112  0.03476 *
## TrD2         0.029688   0.009178   3.235  0.00123 **
## TrD3         0.010264   0.009834   1.044  0.29668
## TrD4         0.004862   0.010072   0.483  0.62929
## TrD5         0.023800   0.010137   2.348  0.01895 *
## TrD6         0.009070   0.009427   0.962  0.33604
## TrD7         0.021698   0.009867   2.199  0.02796 *
## TrD8         0.029825   0.009674   3.083  0.00207 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2466 on 2987 degrees of freedom
## Multiple R-squared:  0.0193, Adjusted R-squared:  0.01667
## F-statistic: 7.348 on 8 and 2987 DF,  p-value: 1.018e-09

confint(m)
```

```
##              2.5 %      97.5 %
## (Intercept)  0.419606599  0.459572237
## TrD1        -0.039576313 -0.001470396
## TrD2         0.011693416  0.047683484
## TrD3        -0.009017722  0.029545928
## TrD4        -0.014885661  0.024610222
## TrD5         0.003923916  0.043676772
## TrD6        -0.009413234  0.027553212
## TrD7         0.002350257  0.041045269
## TrD8         0.010857953  0.048792964
```


From this we can see that the addition of the training data sets TrD 3,4,6 is generally not of relevance for the performance score. Interestingly 4 and 6 also don't appear in the combination of sets with the maximum average test score.

Subset analysis

To analyze the performance that one subset adds to the overall performance we take the subset of all data where a certain training set is used and look if the performance has an observable difference.

```
sub1= subset(d,TrD1==1)
sub2= subset(d,TrD2==1)
sub3= subset(d,TrD3==1)
sub4= subset(d,TrD4==1)
sub5= subset(d,TrD5==1)
sub6= subset(d,TrD6==1)
sub7= subset(d,TrD7==1)
sub8= subset(d,TrD8==1)
boxplot(sub1$score,sub2$score,sub3$score,sub4$score,sub5$score,sub6$score,sub7$score,sub8$score,names =
```

Overall Performance of subsets of all data on a certain training set




From this we can see that the addition of a single training set has no effect on the overall average performance when considering all Test data sets and models.

We visualize the subsets of all combinations of TrD1 and TrD2 and plot their scores against the test scores. Overall, for most training sets the scores are similar. Although for TrD2 there does seem to be a higher score when performed on TeD1, leading to believe that TrD2 performs particularly well on TeD1. For the other Training sets very similar results were observed as for TrD1 thus we choose not to visualize these.


```
plot(sub1$TeD, sub1$score, xlab = "TeD", ylab = "Scores", main= "Performance of subset of all data with T
```

Performance of subset of all data with TrD1 vs TestData set



```
plot(sub2$TeD, sub2$score, xlab = "TeD", ylab = "Scores", main="Performance of subset of all data with TrD1 vs TestData set")
```

Performance of subset of all data with TrD2 vs TestData set



Conclusion

When looking at the combinations of training sets that give the best average performance it can be seen that most of the highest ranking combinations contain the dataset TrD2 and TrD8. This can mean that for good performance on all test sets these test sets are important. We also perform an analysis of the effect of the amount of training sets, this shows that there is indeed a positive effect on the score with an increase in the amount of sets used. Doing an analysis of the performances of the single training set models we can observe that some single training set have a higher average performance in general. TrD 5,7 and 8 seem to provide the best performance. Notably 2, has a very low mean performance but also very high variance. From the analysis of the subset we can see that over all models and combinations of training data sets the performance averages out to the same. From the analysis it looks like the choice of a single model training set does have some effect on the general performance but when averaged out over all combinations there is no effect on the inclusion of a training set.